

DEFORMABLE MEMORY UNDER ACTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study deformable object manipulation. The motivating bet is: Model deformable-object memory as action-conditioned hidden state, not visual appearance.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For deformable object manipulation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2023a; 2024a; 2013; 2023b; 2024b; 2023c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.512	0.488
observed_only_wam	hidden_mode_shift	0.168	0.833
augmentation_as_data_wam	nominal	0.613	0.387
augmentation_as_data_wam	hidden_mode_shift	0.282	0.718
scalar_uncertainty_wam	nominal	0.695	0.305
scalar_uncertainty_wam	hidden_mode_shift	0.343	0.657
proposed_branch_mechanism	nominal	0.792	0.208
proposed_branch_mechanism	hidden_mode_shift	0.557	0.443

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Deformable Memory Under Action argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Visual object detection with deformable part models, 2013. <https://doi.org/10.1145/2494532>.
- Learning language-conditioned deformable object manipulation with graph dynamics, 2023a. <http://arxiv.org/abs/2303.01310v3>.
- Learning foresightful dense visual affordance for deformable object manipulation, 2023b. <http://arxiv.org/abs/2303.11057v3>.
- Sim-to-real transfer for visual reinforcement learning of deformable object manipulation for robot-assisted surgery, 2023c. <https://doi.org/10.1109/lra.2022.3227873>.
- Action-conditional implicit visual dynamics for deformable object manipulation, 2024a. <https://doi.org/10.1177/02783649231191222>.
- Learning dynamic compact memory embedding for deformable visual object tracking, 2024b. <https://doi.org/10.1109/tnnls.2022.3208605>.